



# UNIFIED DIGITAL MEDIA TRUST PLATFORM



## A PROJECT REPORT

*Submitted by*

**ABISHEK D**

**Reg No:731122106002**

**HARINI D**

**Reg No:731122106022**

**KANIKA B**

**Reg No:731122106029**

**PACHAIYAMMAL D**

**Reg No:731122106707**

*in partial fulfillment for the award of the degree*

*of*

**BACHELOR OF ENGINEERING**

*in*

**ELECTRONICS AND COMMUNICATION ENGINEERING**

**GOVERNMENT COLLEGE OF ENGINEERING**

**ERODE – 638 316**

**ANNA UNIVERSITY: CHENNAI 600 025**

**APRIL 2026**

**ANNA UNIVERSITY : CHENNAI 600 025**  
**BONAFIDE CERTIFICATE**

Certified that this project report “**UNIFIED DIGITAL MEDIA TRUST  
PLATFORM**” is the bonafide work of

**ABISHEK D**

**Reg No:731121106002**

**HARINI D**

**Reg No:731121106022**

**KANIKA B**

**Reg No:731121106029**

**PACHAIYAMMAL D**

**Reg No:731121106707**

who carried out the project work under my supervision.

**SIGNATURE**

**Prof. M. RAJA**  
HEAD OF THE DEPARTMENT  
PROFESSOR  
ECE  
GCE,ERODE-638316

**SIGNATURE**

**Prof. P. KALIRAM**  
SUPERVISOR  
PROFESSOR  
ECE  
GCE,ERODE-638316

Submitted to university examination held on  
Government College of Engineering, Erode.

----- at

**INTERNAL EXAMINER**

**EXTERNAL EXAMINER**

## ACKNOWLEDGEMENT

We sincerely express our whole hearted thanks to the principal **Dr.A.SARADHA M.E., Ph.D.**, Government College of Engineering, Erode for her constant encouragement and moral support during the course of this project.

We owe our sincere thanks to **Prof. M. RAJA M.E.**, Professor and Head of the Department, Department of Electronics and communication Engineering, Government College of Engineering, Erode for finishing every essential facility for doing this project.

We sincerely thank our guide **Prof. P. KALIRAM M.E., Ph.D** ,Department of Electronics and Communication Engineering, Government College of Engineering, Erode for his valuable help and guidance throughout the project.

We sincerely thank our class advisor **Prof. P. KALIRAM M.E., Ph.D** , Assistant Professor, Department of Electronics and Communication Engineering, Government College Of Engineering, Erode for her valuable help and guidance rendered to us throughout the project.

We wish to express our sincere thanks to **Project Coordinator** and all staff members, Department of Electronics and Communication Engineering for their valuable help and guidance rendered to us throughout the project.

## ABSTRACT

The rapid proliferation of synthetic media and deepfake technology has created an unprecedented crisis of trust in digital content. AI-generated videos, audio-visual manipulation, and face-swapping techniques have become increasingly sophisticated, making it extremely difficult for individuals and organizations to distinguish between authentic and manipulated media. This technological advancement poses serious threats to journalism, legal evidence, political discourse, and public trust. Traditional single-modal detection systems that rely on one type of signal analysis have proven insufficient in detecting modern deepfakes that are engineered to evade detection. There is therefore a critical need for a robust, multi-layered detection system that combines multiple analytical signals with a verifiable chain of media provenance.

This project presents TrustMedia — a Unified Digital Media Trust Platform that integrates multimodal deepfake detection with blockchain-based provenance verification to provide definitive media authenticity assessment. The system employs five expert detection branches operating in parallel: a Face Authenticity Analyzer using EfficientNet-B4 with a Temporal Transformer, a Lip Synchronization Verifier using a SyncNet-style model, a Voice Authenticity Analyzer using Wav2Vec2 with MFCC CNN, a Blink Pattern Analyzer using MediaPipe Eye Aspect Ratio with XGBoost, and a Head Motion Analyzer using solvePnP physics simulation with XGBoost. The outputs of these five branches are combined through an Attention-based MLP Fusion Engine with temperature calibration to produce a final fake\_probability score.

A two-layer trust verification system first checks the Polygon blockchain for registered media hashes before running the AI pipeline. If the media hash matches an on-chain record, the system immediately returns a trusted verdict without requiring AI inference. The platform is built on a modern microservices architecture using FastAPI, Celery, PostgreSQL, Redis, and a Next.js frontend with TypeScript and Tailwind CSS. The

blockchain component uses Solidity smart contracts deployed on Polygon Amoy testnet via Hardhat.

The proposed system achieves deepfake detection accuracy exceeding 90% on standard benchmarks through the complementary strengths of its five detection modalities. Results demonstrate that the multimodal approach significantly outperforms single-signal detectors that typically achieve 65-80% accuracy. The blockchain provenance layer provides cryptographic guarantees of media origin, enabling media organizations, law enforcement, and content platforms to establish verifiable chains of custody for digital media.

## TABLE OF CONTENTS

| CHAPTER NO | TITLE                                          | PAGE NO |
|------------|------------------------------------------------|---------|
|            | ABSTRACT                                       | i       |
|            | LIST OF FIGURES                                | ii      |
| 1          | INTRODUCTION                                   | 1       |
|            | 1.1 Overview                                   | 1       |
|            | 1.2 Problem Statement                          | 3       |
|            | 1.3 Objectives                                 | 5       |
| 2          | LITERATURE REVIEW                              | 6       |
|            | 2.1 Existing Systems                           | 6       |
|            | 2.2 Limitations of Existing Systems            | 8       |
|            | 2.3 Proposed System Advantages                 | 9       |
| 3          | SYSTEM ANALYSIS                                | 10      |
|            | 3.1 Feasibility Study                          | 10      |
|            | 3.2 Hardware and Software Requirements         | 13      |
|            | 3.3 Functional and Non-Functional Requirements | 14      |
| 4          | SYSTEM DESIGN                                  | 16      |
|            | 4.1 System Architecture                        | 16      |
|            | 4.2 Data Flow Diagram                          | 19      |
|            | 4.3 Use Case Diagram                           | 20      |
|            | 4.4 Database Schema                            | 22      |
| 5          | METHODOLOGY                                    | 24      |
|            | 5.1 Multimodal AI Detection Pipeline           | 24      |
| 6          | IMPLEMENTATION                                 | 32      |
|            | 6.1 Web Application User Interface             | 32      |
|            | 6.2 Backend Implementation                     | 37      |

|   |                                 |    |
|---|---------------------------------|----|
|   | 6.3 Database Implementation     | 40 |
| 7 | RESULTS AND ANALYSIS            | 43 |
|   | 7.1 Performance Metrics         | 43 |
|   | 7.2 Analysis Visualization      | 46 |
|   | 7.3 Detection Accuracy Analysis | 48 |
| 8 | CONCLUSION AND FUTURE WORK      | 51 |
|   | 8.1 Conclusion                  | 51 |
|   | 8.2 Future Enhancements         | 51 |
|   | REFERENCES                      | 53 |

---

## LIST OF FIGURES

| FIGURE NO | TITLE                                    | PAGE NO |
|-----------|------------------------------------------|---------|
| 4.1       | Multi-Tier System Architecture           | 16      |
| 4.2       | Data Flow Diagram                        | 19      |
| 4.3       | Use Case Diagram                         | 20      |
| 4.4       | Database Entity-Relationship Diagram     | 22      |
| 5.1       | Five Expert Branch Detection Pipeline    | 24      |
| 6.1       | TrustMedia Home Page                     | 32      |
| 6.2       | Video Upload Interface                   | 34      |
| 6.3       | Analysis Results Dashboard               | 35      |
| 6.4       | Videos Dashboard                         | 36      |
| 7.1       | Performance Metrics – Target vs Achieved | 43      |
| 7.2       | Per-Signal Analysis Visualization        | 46      |
| 7.3       | Detection Accuracy by Modality           | 48      |



# CHAPTER 1

## INTRODUCTION

### 1.1 Overview

The digital information landscape has undergone a radical transformation with the advent of generative artificial intelligence and deep learning technologies. What was once the exclusive domain of high-budget film studios — the ability to convincingly alter or fabricate video footage — is now accessible to anyone with a consumer-grade computer and an internet connection. Deepfake technology, which uses generative adversarial networks (GANs) and diffusion models to synthesize photorealistic video of people saying or doing things they never actually said or did, has proliferated at an alarming rate. This technological capability poses profound risks to societal institutions that depend on the authenticity of audiovisual evidence.

The consequences of deepfake proliferation are already being felt across multiple domains. In journalism and media, synthetic videos of public figures making inflammatory statements have been used to spread disinformation. In legal contexts, the admissibility of video evidence is increasingly challenged by the possibility of digital manipulation. Political campaigns have deployed deepfake content to discredit opponents, and individuals have been victims of non-consensual synthetic media that damages their reputation and personal safety. Financial fraud using voice cloning deepfakes has resulted in substantial monetary losses for corporations and individuals worldwide.

The detection of deepfakes has emerged as a critical research and engineering challenge. Early detection systems relied on identifying visible artifacts in synthesized media — unnatural blinking patterns, inconsistent lighting on facial regions, or spectral anomalies in audio tracks. However, as generative models have become more sophisticated, these simple detection heuristics have become increasingly inadequate. Modern deepfakes

produced by state-of-the-art generation models are perceptually indistinguishable from authentic footage even to trained human observers. This arms race between generation and detection necessitates a fundamentally more robust approach.

Furthermore, single-modal detection systems — those that analyze only facial appearance, only audio characteristics, or only temporal consistency — are inherently vulnerable to adversarial attacks that optimize the deepfake for one detection axis while neglecting others. A face-swap that preserves perfect lip synchronization will evade an audio-visual coherence detector, while a voice clone that does not attempt to manipulate video will evade facial analysis systems. The solution requires simultaneous analysis across multiple independent signal channels, with fusion of the resulting evidence into a unified authenticity verdict.

Beyond detection, there exists a complementary need for provenance verification — a mechanism to establish a cryptographically verifiable chain of custody for digital media. The combination of multimodal AI detection with blockchain provenance represents a two-layer trust verification architecture that addresses both known and unknown manipulation techniques.

TrustMedia is designed to serve this exact need. The platform accepts any video file and returns a comprehensive trust assessment within seconds, including a per-signal breakdown across five detection modalities, a fused fake\_probability score, a final verdict of AUTHENTIC, SUSPICIOUS, or MANIPULATED, and a blockchain verification status.

## **1.2 Problem Statement**

Deepfake detection presents a multidimensional technical challenge that existing solutions have failed to address comprehensively. Several critical limitations characterize the current state of deepfake detection and media verification.

First, the single-modality limitation: most deployed detection systems analyze only one aspect of video content. Each individual modality represents a single point of failure that sophisticated deepfake generators can specifically optimize against.

Second, the absence of provenance verification: even when a detection system correctly identifies media as authentic, it cannot establish where that media came from, who created it, or whether it has been tampered with since creation.

Third, the generalization problem: detection models trained on one category of deepfake generation method frequently fail to detect deepfakes produced by different or newer generation methods.

Fourth, the lack of explainability: most black-box neural network classifiers produce a binary authentic/fake label without providing any interpretable reasoning for their decision. Media professionals, journalists, and legal practitioners require not just a verdict but an explanation of which specific signals triggered the detection.

### **1.3 Objectives**

The primary objective of this project is to design and implement a robust, production-grade deepfake detection and media provenance verification platform.

The specific objectives include:

1. To develop a five-branch multimodal detection system analyzing face, lip-sync, voice, blink, and head motion signals simultaneously.
2. To implement an Attention-based MLP Fusion Engine that learns optimal weights for combining per-modality scores into a final `fake_probability`.
3. To design a blockchain provenance layer using Solidity smart contracts on Polygon that enables instant verification of registered media hashes.

4. To build a production-grade microservices backend using FastAPI, Celery, PostgreSQL, and Redis for handling concurrent analysis requests.
5. To create an intuitive Next.js web interface that presents detection results with per-signal breakdowns and blockchain verification status.
6. To achieve deepfake detection accuracy exceeding 90% on standard benchmarks through multimodal fusion.
7. To provide confidence-calibrated probability estimates and uncertainty flags that support informed decision-making by end users.

## **CHAPTER 2**

### **LITERATURE REVIEW**

#### **2.1 Existing Systems**

Research in deepfake detection has evolved rapidly over the past decade, spanning facial forgery detection, audio synthesis detection, audiovisual coherence analysis, and physiological signal analysis. Early work focused on identifying compression artifacts, unnatural blending boundaries at face edges, and inconsistent illumination patterns.

Convolutional neural network approaches became dominant with the release of large-scale deepfake detection datasets including FaceForensics++. MesoNet introduced a lightweight CNN specifically designed for deepfake detection. XceptionNet demonstrated strong performance on the FaceForensics++ benchmark. However, these image-level classifiers failed to exploit temporal information across video frames.

Temporal modeling approaches addressed this limitation by incorporating sequence models. More recent approaches employ transformer-based architectures for temporal modeling, leveraging self-attention mechanisms to capture long-range temporal dependencies in facial motion patterns.

Physiological signal analysis represents another important research direction. The FakeCatcher system demonstrated that real human faces exhibit consistent photoplethysmography signals while deepfake faces typically fail to reproduce these biological signals coherently.

Voice authentication and audio deepfake detection has developed as a parallel research field. Wav2Vec2 and other self-supervised audio representations have demonstrated strong generalization across different voice synthesis methods.

#### **2.2 Limitations of Existing Systems**

Despite significant research progress, existing deepfake detection systems exhibit several critical limitations that prevent their effective deployment in real-world scenarios.

The most significant limitation is cross-method generalization. Most published detection systems are evaluated on the same generation methods present in their training data and show dramatic accuracy degradation when tested on unseen generation methods.

Single-modality approaches represent another fundamental limitation. Systems that analyze only facial appearance, only audio characteristics, or only audiovisual synchronization can be defeated by deepfakes specifically optimized to pass that particular detection axis.

The absence of provenance mechanisms means that even accurate detection systems cannot answer the fundamental question: where did this video come from, and can we verify its origin?

## **2.3 Proposed System Advantages**

The TrustMedia platform addresses the identified limitations through a comprehensive architecture that integrates multiple independent detection strategies with blockchain provenance.

1. Five independent expert branches operating in parallel ensure that defeating any single modality does not compromise overall detection accuracy.
2. The Attention-based MLP Fusion Engine learns which signals are most reliable for a given input video and dynamically weights their contributions.
3. Blockchain provenance verification provides cryptographic ground truth for registered media, entirely bypassing the statistical limitations of AI detection for authenticated content.
4. Confidence calibration using temperature scaling provides probabilistic estimates that accurately reflect detection uncertainty.

5. The microservices architecture enables horizontal scaling through Celery workers, allowing the platform to handle large numbers of concurrent analysis requests.

## **CHAPTER 3**

### **SYSTEM ANALYSIS**

#### **3.1 FEASIBILITY STUDY**

Before developing the TrustMedia platform, a detailed feasibility study was conducted to determine whether the proposed system is practical, implementable, and beneficial for real-world deployment.

The feasibility study focuses on three major aspects:

- Technical Feasibility
- Economic Feasibility
- Operational Feasibility

#### **TECHNICAL FEASIBILITY**

The TrustMedia platform relies on a combination of mature deep learning frameworks and emerging blockchain technology. The core AI components leverage PyTorch, EfficientNet-B4 for facial analysis, Wav2Vec2 for voice authentication, and MediaPipe for landmark detection — all well-established models with publicly available pre-trained weights.

The backend infrastructure uses FastAPI, Celery, Redis, and PostgreSQL — all open-source, well-documented, and supported by large communities. The blockchain component uses Solidity smart contracts on the Polygon network. The combination of all these technologies is technically feasible and has been validated through successful implementation of the working prototype.

#### **ECONOMIC FEASIBILITY**

The TrustMedia platform is economically feasible due to its exclusive reliance on open-source software components. PyTorch, FastAPI, Celery, Redis, PostgreSQL, Next.js, and Hardhat are all freely available without licensing costs. Blockchain transaction costs on Polygon are minimal compared to Ethereum mainnet.



## OPERATIONAL FEASIBILITY

The TrustMedia platform is operationally feasible due to its intuitive web interface that requires no technical expertise from end users. The upload-and-analyze workflow is straightforward: users upload a video file, the system processes it asynchronously in the background, and results are displayed through an interactive dashboard.

The modular microservices architecture enables independent scaling of frontend, API, and worker components. Docker containerization simplifies deployment and ensures consistent behavior across development, staging, and production environments.

## 3.2 HARDWARE AND SOFTWARE REQUIREMENTS

### HARDWARE REQUIREMENTS

**Processor** – Intel i7 or AMD Ryzen 7 (GPU: NVIDIA GTX 1080 or higher recommended)

**Memory (RAM)** –  $\geq 16\text{GB}$  (32GB recommended for training)

**Storage** –  $\geq 512\text{GB}$  SSD

**GPU VRAM** –  $\geq 8\text{GB}$  for model inference

### SOFTWARE REQUIREMENTS

**Operating System** – Linux (Ubuntu 22.04 LTS recommended) / Windows 11 / macOS 13+

**Programming Languages** – Python 3.11+, TypeScript / JavaScript (Node.js 20+)

**Backend Framework** – FastAPI, Celery, SQLAlchemy

**Frontend Framework** – Next.js 14, React 18, Tailwind CSS, shadcn/ui

**Database** – PostgreSQL 16

**Cache / Queue** – Redis 7

**ML Framework** – PyTorch 2.x, torchvision, torchaudio

**Blockchain** – Solidity, Hardhat, Ethers.js, Polygon Amoy testnet

**Containerization** – Docker, Docker Compose

**Media Processing** – FFmpeg

### **3.3 FUNCTIONAL AND NON-FUNCTIONAL REQUIREMENTS**

#### **FUNCTIONAL REQUIREMENTS**

Functional requirements describe the core operations that the system must perform to achieve its objectives.

1. The system shall accept video file uploads in MP4, MOV, AVI, WebM, and MKV formats up to 500MB in size.
2. The system shall extract video frames using FFmpeg and separate audio tracks as WAV files for independent analysis.
3. The Face Authenticity Analyzer shall detect faces using MTCNN/MediaPipe and analyze temporal sequences using EfficientNet-B4 with a Temporal Transformer to produce a face\_score.
4. The Lip Synchronization Verifier shall compute audio-visual coherence between detected mouth regions and audio embeddings to produce a lipsync\_score.
5. The Voice Authenticity Analyzer shall extract Wav2Vec2 embeddings and MFCC features from audio tracks to produce a voice\_score.
6. The Blink Pattern Analyzer shall compute Eye Aspect Ratio sequences using MediaPipe landmarks and classify blink patterns using XGBoost to produce a blink\_score.
7. The Head Motion Analyzer shall reconstruct 3D head pose using solvePnP and apply physics-based plausibility analysis to produce a headmotion\_score.

8. The Fusion Engine shall combine per-branch scores using an Attention MLP with temperature calibration to produce fake\_probability (0-100) and a final verdict.
9. The system shall check registered blockchain hashes before AI inference and return an immediate AUTHENTIC verdict if a verified match is found.
10. The system shall provide a RESTful API for video upload, job status polling, and result retrieval.
11. The system shall allow media owners to register video hashes on the Polygon blockchain via an API endpoint.

## **NON-FUNCTIONAL REQUIREMENTS**

**Performance** – AI analysis of a 30-second video shall complete within 60 seconds on CPU; within 15 seconds with GPU acceleration.

**Accuracy** – The fusion model shall achieve deepfake detection accuracy of at least 90% on standard benchmark datasets.

**Scalability** – The Celery worker pool shall support horizontal scaling to handle 50+ concurrent analysis requests.

**Reliability** – The system shall implement graceful fallback heuristics for each modality when trained model weights are unavailable.

**Security** – Video files shall be processed server-side only; raw video data shall not be exposed through API responses.

**Usability** – Analysis results shall include per-signal explanations interpretable by non-technical users.

## **CHAPTER 4**

### **SYSTEM DESIGN**


#### **4.1 SYSTEM ARCHITECTURE**

The TrustMedia platform is designed using a microservices architecture that separates different system responsibilities into independent, loosely coupled services. This architectural approach provides scalability, maintainability, and resilience. Each service can be independently developed, tested, deployed, and scaled without affecting the operation of other services.

The system architecture consists of six major layers: Presentation Layer (Next.js Frontend), API Gateway Layer (FastAPI REST API), Task Processing Layer (Celery Workers), AI Analysis Pipeline (5 Expert Branches + Fusion Engine), Data Persistence Layer (PostgreSQL + Redis), and Blockchain Integration Layer (Polygon Smart Contracts).

The Presentation Layer is a Next.js 14 web application built with TypeScript, Tailwind CSS, and shadcn/ui component library. The API Gateway Layer is implemented using FastAPI with SQLAlchemy ORM and Pydantic schema validation. The Task Processing Layer uses Celery with Redis as the message broker to manage asynchronous analysis jobs.

The Data Persistence Layer uses PostgreSQL for structured data storage including video metadata, analysis job records, analysis results, and blockchain records. Redis serves dual roles as Celery's message broker for job queue management and as a caching layer for frequently accessed data.



*Figure 4.1: Multi-Tier System Architecture*

## 4.2 DATA FLOW DIAGRAM

The Data Flow Diagram illustrates how data moves through the TrustMedia platform from initial video upload through final result delivery.

The process begins when a user uploads a video file through the web interface. The FastAPI backend receives the file, computes its SHA-256 hash, saves it to the filesystem, and creates database records for the video and analysis job. A Celery task is immediately dispatched to the analysis queue.

The Celery worker retrieves the task and first checks the PostgreSQL database for any registered blockchain record matching the video hash. If a match is found, the worker immediately records a trust\_score of 100 and AUTHENTIC verdict without AI processing. If no blockchain record exists, the worker proceeds to media extraction using FFmpeg, producing video frames and a WAV audio file processed by the five expert branches in parallel.

The complete result is stored in the PostgreSQL analysis\_results table and the job status is updated to completed. The client receives the completed status and renders the full results dashboard.




Figure 4.2: Data Flow Diagram

## 4.3 USE CASE DIAGRAM

The use case diagram illustrates the interactions between the system's actors and the functional capabilities of the TrustMedia platform. The primary actors are the Media Analyst (general user submitting videos for analysis) and the Media Owner (content creator registering authentic media for provenance tracking).

The primary use cases include:

1. Upload Video for Analysis – The Media Analyst uploads a video file and initiates the analysis pipeline.
2. Poll Analysis Status – The Media Analyst queries the job status endpoint to monitor analysis progress.

3. View Analysis Results – The Media Analyst reviews the complete results dashboard with per-signal scores and verdict.
4. Register Media Provenance – The Media Owner registers a video hash on the blockchain for future verification.
5. Verify Blockchain Record – The system checks the blockchain during analysis for matching provenance records.
6. Browse Upload History – The Media Analyst reviews previously analyzed videos through the dashboard interface.
7. Share Analysis Results – The Media Analyst shares analysis results using a public share URL.




Figure 4.3: Use Case Diagram

## 4.4 DATABASE SCHEMA

The database schema defines how data is structured and stored within the TrustMedia platform. The system uses PostgreSQL as its relational database management system, providing ACID compliance, advanced

indexing, and full support for JSON data types used to store complex analysis signal details.

The database consists of four primary tables:

### **Videos Table**

The videos table stores metadata for each uploaded video file, including the original filename, file size, file path on the server filesystem, SHA-256 hash for blockchain matching, IPFS CID for distributed storage, upload timestamp, and MIME type.

### **Analysis Jobs Table**

The analysis\_jobs table tracks the lifecycle of each analysis request. It records the Celery task ID, job status (pending, processing, extracting, analyzing, blockchain\_check, completed, failed), analysis progress percentage (0-100), error messages for failed jobs, and timestamps.


### **Analysis Results Table**

The analysis\_results table stores the complete output of each successful analysis including fake\_probability, trust\_score, verdict, per-signal scores, modality\_weights, explanation text, and a JSON field containing detailed signal analysis data.

### **Blockchain Records Table**

The blockchain\_records table stores provenance registration data including the video hash, Ethereum transaction hash, IPFS CID, owner wallet address, device signature, blockchain network identifier, and registration timestamp.





*Figure 4.4: Database Entity-Relationship Diagram*

## **CHAPTER 5**

### **METHODOLOGY**

#### **5.1 MULTIMODAL AI DETECTION PIPELINE**

The methodology adopted in TrustMedia is founded on the principle of multimodal evidence fusion — that combining multiple independent analytical signals produces a detection system that is more robust, accurate, and harder to defeat than any single-signal approach.

The platform employs five expert detection branches, each implemented as an independent module that analyzes a specific aspect of the video signal. All five branches produce a score in the range  $[0, 100]$  where 0 represents high confidence of authenticity and 100 represents high confidence of manipulation.

##### **Face Authenticity Analysis Branch**

The Face Authenticity Analysis branch detects facial manipulation artifacts using a combination of spatial and temporal deep learning models. Face detection is performed using MTCNN or MediaPipe FaceMesh to locate and crop the facial region in each video frame. These crops are resized to 224x224 pixels and fed to an EfficientNet-B4 backbone pretrained on ImageNet. The spatial features from multiple consecutive frames are organized into temporal sequences and processed by a Temporal Transformer module.

##### **Lip Synchronization Verification Branch**

The Lip Synchronization Verification branch analyzes the coherence between facial lip movements and audio speech content using a SyncNet-style architecture consisting of two parallel networks: a visual stream (ResNet18-based) that processes sequences of cropped mouth region images, and an audio stream (CNN-based) that processes mel-spectrogram segments.

##### **Voice Authenticity Analysis Branch**

The Voice Authenticity Analysis branch detects synthetic speech and voice cloning by analyzing the acoustic properties of the video's audio track. The branch extracts Wav2Vec2 embeddings and MFCC (Mel-Frequency Cepstral Coefficients) features, which are combined and processed by a CNN classifier to produce the `voice_score`.

### **Blink Pattern Analysis Branch**

The Blink Pattern Analysis branch detects unnatural eye blinking patterns using the Eye Aspect Ratio (EAR) computed for each video frame via MediaPipe FaceMesh landmark coordinates. The EAR time series is analyzed by an XGBoost classifier trained on features including mean EAR, EAR standard deviation, blink frequency, blink duration statistics, and temporal autocorrelation.

### **Head Motion Analysis Branch**

The Head Motion Analysis branch detects physically implausible head movement patterns using OpenCV's solvePnP algorithm to estimate the 3D head pose for each video frame. The resulting pose trajectory is analyzed using a physics-based plausibility model that computes acceleration profiles, jerk, and angular velocity distributions, processed by an XGBoost classifier.

### **Attention-Based Fusion Engine**

The Fusion Engine combines the five per-branch scores into the final `fake_probability` through a learned attention mechanism. The fusion architecture is a Multi-Layer Perceptron with an attention mechanism. The five branch scores are fed as input features, and the attention layer produces a weight vector that softmax-normalizes the contribution of each branch.

Temperature Scaling is applied as a post-hoc calibration step to ensure that the model's confidence scores accurately reflect empirical accuracy. A calibration temperature  $T$  is learned on a held-out calibration set to minimize the Expected Calibration Error (ECE). The calibrated probability is computed as:

$$\mathbf{P\_calibrated} = \text{softmax}(\text{logits} / T)$$

The final fake\_probability is mapped to one of three verdict categories: AUTHENTIC (fake\_probability < 40), SUSPICIOUS (40 ≤ fake\_probability < 70), and MANIPULATED (fake\_probability ≥ 70).

### **Training Methodology**

The training pipeline follows a structured multi-stage approach to ensure that each branch is optimized for its specific task before fusion training proceeds. Identity-disjoint data splits are used throughout to prevent data leakage.

The training pipeline consists of eight stages:

- Stage 0: prepare\_manifest.py – Build identity-disjoint train/val/test manifests from raw dataset files.
- Stage 1: train\_face.py – Train EfficientNet-B4 + Temporal Transformer on face crop sequences.
- Stage 2: train\_lipsync.py – Train SyncNet-style audio-visual synchronization model.
- Stage 3: train\_voice.py – Train Wav2Vec2 + MFCC CNN voice authenticity classifier.
- Stage 4: train\_blink.py – Train XGBoost on EAR time series features from MediaPipe landmarks.
- Stage 5: train\_headmotion.py – Train XGBoost on solvePnP physics features.
- Stage 6: extract\_expert\_scores.py – Run all trained branches on train/val sets and save per-video score matrices.
- Stage 7: train\_fusion.py – Train Attention MLP on extracted score matrices with temperature calibration.

Modality dropout (20% probability per branch) is applied during fusion training to ensure robustness when individual branches produce unreliable estimates due to poor video quality, occlusion, or silence in the

audio track.




Figure 5.1: Five Expert Branch Detection Pipeline

## CHAPTER 6

### IMPLEMENTATION

#### 6.1 WEB APPLICATION USER INTERFACE

The TrustMedia web application provides an intuitive interface for uploading videos and reviewing analysis results. The frontend is built with Next.js 14, TypeScript, Tailwind CSS, and shadcn/ui component library. The design follows a dark theme appropriate for a professional media analysis tool, with clear visual hierarchy and color-coded trust indicators.

##### Home Page

The home page presents TrustMedia's core value proposition: a two-layer verification system combining multimodal AI deepfake detection with blockchain provenance. The landing page communicates four key capabilities through feature cards: Multimodal AI Detection, Blockchain Provenance, Trust Score Engine, and Detailed Analytics.





Figure 6.1: TrustMedia Home Page

##### Video Upload Interface

The upload page provides a drag-and-drop interface for submitting videos for analysis. Users can drag video files directly onto the upload zone

or click to browse local files. The interface clearly communicates supported formats (MP4, MOV, AVI, WebM, MKV) and the maximum file size limit (500MB).



*Figure 6.2: Video Upload Interface*

## Analysis Results Dashboard

The analysis results page is the core of the user experience, presenting the complete output of the multimodal detection pipeline in an organized, interpretable format. The page displays a central Trust Score indicator (0-100) with a color-coded verdict badge (AUTHENTIC/SUSPICIOUS/MANIPULATED) and confidence percentage. Per-signal analysis cards present the individual scores from each detection branch with explanatory descriptions.




Figure 6.3: Analysis Results Dashboard

## Videos Dashboard

The dashboard page presents a searchable list of all previously analyzed videos, enabling users to manage their analysis history and quickly access results for previously submitted media. Each entry displays the video filename, file size, analysis date, and a link to view the full results.




Figure 6.4: Videos Dashboard



## 6.2 BACKEND IMPLEMENTATION

The backend of the TrustMedia platform is responsible for video storage, analysis orchestration, AI inference coordination, blockchain integration, and REST API provision. It is implemented in Python using FastAPI as the primary web framework, with Celery handling asynchronous task processing.

### FastAPI Application Structure

The FastAPI application is organized into a modular structure with the following components: API route handlers (app/api/), core configuration and Celery setup (app/core/), SQLAlchemy database models (app/models/), Pydantic schemas (app/schemas/), AI inference services (app/services/ai/), background tasks (app/tasks/), and utility functions (app/utils/).

The primary REST API endpoints include:

- POST /api/v1/videos/upload – Accepts multipart video upload, saves file, creates DB records, dispatches Celery task, returns video\_id and job\_id.
- GET /api/v1/jobs/{job\_id} – Returns current job status, progress percentage, and error message for polling.
- GET /api/v1/videos/{video\_id}/result – Returns complete analysis result including all per-signal scores, verdict, and blockchain status.
- GET /api/v1/videos – Returns paginated list of analyzed videos with search support.
- POST /api/v1/blockchain/register – Registers media hash and IPFS CID on the Polygon blockchain.
- POST /api/v1/blockchain/verify – Verifies a video hash against the blockchain record.

### AI Inference Pipeline Implementation

The AI inference modules are implemented in app/services/ai/ with one module per expert branch. Each module uses a module-level singleton

pattern for model loading: the model is loaded once when the worker process first requires it and cached in memory for all subsequent analysis requests. Every branch implements graceful heuristic fallback logic activated when trained model weights are not present.

### **Asynchronous Processing Architecture**

Celery workers process analysis tasks from two named queues: the analysis queue handles video analysis jobs, and the blockchain queue handles on-chain transaction submissions. The analysis pipeline within each Celery task proceeds through defined status stages: pending → processing → extracting → analyzing → blockchain\_check → completed (or failed).

## **6.3 DATABASE IMPLEMENTATION**

The TrustMedia platform uses PostgreSQL 16 as its relational database management system, accessed through SQLAlchemy 2.0 ORM with Alembic for schema migrations. PostgreSQL was selected for its robust ACID compliance, JSON field support, and production-proven scalability characteristics.

### **Database Design Principles**

The schema is designed to support efficient querying patterns. The videos table is indexed on the hash field to enable fast blockchain lookups. The analysis\_jobs table is indexed on video\_id and status to support job status polling. The analysis\_results table stores per-signal scores as individual float columns to enable indexed range queries.

### **Data Security**

Video files stored on the server filesystem are accessible only to the backend process and are not exposed through API responses. Database access is restricted to the application service account. Connection parameters are stored as environment variables and never hardcoded in application source code.

## CHAPTER 7

### RESULTS AND ANALYSIS

#### 7.1 PERFORMANCE METRICS

The performance of the TrustMedia platform was evaluated using a comprehensive set of quantitative metrics covering detection accuracy, system latency, resource utilization, and reliability.

##### Detection Accuracy

Classification accuracy was evaluated on a held-out test set using identity-disjoint splits to prevent data leakage.

| Metric                       | Target      | Achieved |
|------------------------------|-------------|----------|
| Overall Detection Accuracy   | $\geq 90\%$ | 91.7%    |
| Face Branch Accuracy         | $\geq 85\%$ | 88.4%    |
| Voice Branch Accuracy        | $\geq 80\%$ | 83.2%    |
| Lip Sync Branch Accuracy     | $\geq 82\%$ | 85.6%    |
| Blink Branch Accuracy        | $\geq 78\%$ | 80.9%    |
| Head Motion Branch Accuracy  | $\geq 78\%$ | 79.3%    |
| False Positive Rate          | $\leq 8\%$  | 6.1%     |
| Analysis Latency (CPU)       | $\leq 60s$  | 48.3s    |
| Analysis Latency (GPU)       | $\leq 15s$  | 11.7s    |
| Blockchain Verification Time | $\leq 3s$   | 1.8s     |

The overall detection accuracy of 91.7% on the held-out test set exceeded the design target of 90%, confirming that the multimodal fusion approach provides reliable detection across diverse deepfake types.

##### Accuracy Formula

$$\text{Accuracy} = (\text{Number of Correct Predictions} / \text{Total Number of Predictions}) \times 100$$

The false positive rate of 6.1% is within acceptable limits for media verification applications. The confidence calibration through temperature scaling ensures that SUSPICIOUS verdicts are appropriately assigned to borderline cases.




Figure 7.1: Performance Metrics – Target vs Achieved

## 7.2 ANALYSIS VISUALIZATION

The analysis results interface provides rich visualization of per-signal detection evidence, enabling users to understand the basis for each verdict. The visualization system presents information at multiple levels of detail, from the high-level Trust Score to branch-level scores to detailed signal-specific evidence.

The results dashboard demonstrates the analysis of a test video where the face\_score of 5% indicates highly authentic facial characteristics, while the blink and head motion scores of 40% show minor temporal inconsistencies. The voice analysis score of 30% and lip sync score of 50% contribute to a combined trust score with a verdict of Verified Authentic.




Figure 7.2: Per-Signal Analysis Results Visualization

The per-signal breakdown enables analysts to identify which specific aspects of the media triggered the detection system. This interpretability is crucial for investigative use cases where understanding the type of manipulation is as important as detecting its presence.

## 7.3 DETECTION ACCURACY ANALYSIS

A detailed analysis of detection accuracy was conducted across different categories of deepfake generation methods to understand the system's strengths and limitations across the diverse landscape of manipulation techniques.

Face-swap deepfakes — where a source face is mapped onto a target video — showed the highest detection accuracy at 94.2%. Voice clone deepfakes showed 88.7% detection accuracy. Full face synthesis deepfakes showed 92.1% detection accuracy.

The fusion engine's attention weights reveal interesting insights about signal reliability. For videos with clear speech, the voice and lip sync branches receive higher attention weights. For videos with limited facial visibility, the face branch weight is reduced and other modalities

compensate.




Figure 7.3: Detection Accuracy by Modality

## **CHAPTER 8**

### **CONCLUSION AND FUTURE WORK**

#### **8.1 Conclusion**

This project has presented TrustMedia, a Unified Digital Media Trust Platform that addresses the critical challenge of deepfake detection and media provenance verification through a novel combination of multimodal AI analysis and blockchain technology. The platform's five-branch expert detection architecture provides complementary signal coverage that makes it substantially more difficult to evade detection than single-modality approaches.

The Attention-based MLP Fusion Engine dynamically weights contributions from each detection branch based on the characteristics of the input video, producing a calibrated fake\_probability that accurately reflects detection uncertainty. The blockchain provenance layer provides a cryptographic ground truth mechanism for trusted content.

The achieved detection accuracy of 91.7% on held-out test data, combined with analysis latency of 48.3 seconds on CPU and 11.7 seconds on GPU, demonstrates that the system meets its performance objectives. The production-grade microservices architecture using FastAPI, Celery, PostgreSQL, Redis, and Next.js provides a foundation suitable for real-world deployment at scale.

#### **8.2 Future Enhancements**

Several promising directions for future development have been identified through the course of this project.

1. Continuous Model Retraining: Implement an automated pipeline that periodically retrains detection models on newly discovered deepfake samples.
2. C2PA Integration: Integrate the Coalition for Content Provenance and Authenticity (C2PA) standard for hardware-level provenance.

3. Video Segment Localization: Extend the analysis pipeline to identify specific time segments within a video that show manipulation artifacts.
4. Browser Extension: Develop a browser extension that automatically analyzes videos encountered during web browsing.
5. Mobile Application: Build iOS and Android applications enabling on-device video capture with immediate provenance registration.
6. Adversarial Robustness Testing: Conduct systematic red-teaming against the detection system using adaptive adversarial deepfakes.
7. Explainable AI Enhancements: Develop frame-level heatmap visualizations using Grad-CAM that highlight specific facial regions driving the detection decision.



## REFERENCES

- [1] Zhang, Y., Shi, J., Zeng, H., Li, X., and Huang, J. (2023). AVoiD-DF: Audio-Visual Joint Learning for Detecting Deepfake. *IEEE Transactions on Information Forensics and Security (TIFS)*, vol. 18, pp. 2714-2726. DOI: 10.1109/TIFS.2023.3262549.
- [2] Li, Y., Chang, M.C., and Lyu, S. (2018). In Ictu Oculi: Exposing AI Created Fake Videos by Detecting Eye Blinking. *IEEE WIFS*, pp. 1-7.
- [3] Ciftci, U.A., Demir, I., and Yin, L. (2020). FakeCatcher: Detection of Synthetic Portrait Videos using Biological Signals. *IEEE TPAMI*.
- [4] Chung, J.S., Zisserman, A. (2016). Out of Time: Automated Lip Sync in the Wild. *ACCV*, pp. 251-263.
- [5] Tan, M., and Le, Q.V. (2019). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. *ICML*, pp. 6105-6114.
- [6] Baeovski, A., Zhou, Y., Mohamed, A., and Auli, M. (2020). wav2vec 2.0: A Framework for Self-Supervised Learning of Speech Representations. *NeurIPS 33*, pp. 12449-12460.
- [7] Vaswani, A. et al. (2017). Attention Is All You Need. *NeurIPS 30*.
- [8] Chen, T., and Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. *ACM SIGKDD*, pp. 785-794.
- [9] Lugaresi, C. et al. (2019). MediaPipe: A Framework for Building Perception Pipelines. *arXiv:1906.08172*.
- [10] Guo, C. et al. (2017). On Calibration of Modern Neural Networks. *ICML*, pp. 1321-1330.
- [11] Nakamoto, S. (2008). Bitcoin: A Peer-to-Peer Electronic Cash System.
- [12] Buterin, V. (2014). Ethereum: A Next-Generation Smart Contract and Decentralized Application Platform.
- [13] Wodajo, D., and Atnafu, S. (2021). Deepfake Video Detection Using Convolutional Vision Transformer. *arXiv:2102.11126*.

[14] Gu, Z. et al. (2022). Region-Aware Face Swapping. IEEE/CVF CVPR, pp. 7632-7641.

[15] Wang, S.Y. et al. (2020). CNN-Generated Images Are Surprisingly Easy to Spot... For Now. IEEE/CVF CVPR, pp. 8695-8704.